

Lecture 7: Linear-Time Sorting

Linear time sorting is possible sometimes. This lecture shows when


Lecture Overview

- Comparison model
- Lower bounds
 - searching: $\Omega(\lg n)$
 - sorting: $\Omega(n \lg n)$
- $O(n)$ sorting algorithms for small integers
 - counting sort
 - radix sort

Lower Bounds

Claim

- searching among n preprocessed items requires $\Omega(\lg n)$ time
 \implies binary search, AVL tree search optimal
- sorting n items requires $\Omega(n \lg n)$
 \implies mergesort, heap sort, AVL sort optimal



...in the comparison model

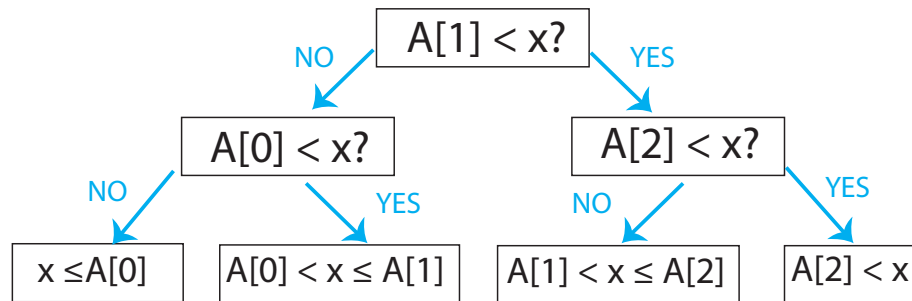
Comparison Model of Computation

- input items are black boxes (ADTs)
- only support comparisons ($<$, $>$, \leq , etc.)
- time cost = # comparisons

Decision Tree

Any comparison algorithm can be viewed/specified as a tree of all possible comparison outcomes & resulting output, for a particular n :

- example, binary search for $n = 3$:



- internal node = binary decision
- leaf = output (algorithm is done)
- root-to-leaf path = algorithm execution
- path length (depth) = running time
- height of tree = worst-case running time

You don't normally write a binary search as tree
But its useful to show via tree
"this is all this possibilities this algorithm can do"

In fact, binary decision tree model is more powerful than comparison model, and lower bounds extend to it

Search Lower Bound

Assumes you have n preprocessed items initially. Do anything with them at $t=0$
No matter what you do, to find a new item, it will take $\log n$ time

- # leaves \geq # possible answers $\geq n$ (at least 1 per $A[i]$)
 - decision tree is binary And it must have n leaves, 1 for each answer. Maybe more than 1 leaf correspond to 1 answer. (multiple paths to same answer). But leaves have to be atleast n
 - \Rightarrow height $\geq \lg \Theta(n) = \lg n \pm \Theta(1)$ This extends to finding successor, predecessor etc. All of them must take $\log n$ time. Searching takes $\log n$ time. Hence BST are very efficient (Yes you can find pred/succ in $O(1)$ time in sorted arrays, but search is still $O(\log n)$ sorted array insertion is $O(n)$ time)
- And height is the worst case running time

Sorting Lower Bound

- leaf specifies answer as permutation: $A[3] \leq A[1] \leq A[9] \leq \dots$ $\log n$ is the best you can do in comparison model, and BST archives that
- all $n!$ are possible answers

Just like case of search;
decision tree is binary; #leaves \geq #answers
#leaves $\geq n!$

- # leaves $\geq n!$ height is $\log(\text{leaves})$ in binary tree

$$\begin{aligned}
 \Rightarrow \text{height} &\geq \lg n! \\
 &= \lg(1 \cdot 2 \cdots (n-1) \cdot n) \\
 &= \lg 1 + \lg 2 + \cdots + \lg(n-1) + \lg n \\
 &= \sum_{i=1}^n \lg i \\
 &\geq \sum_{i=n/2}^n \lg i \\
 &\geq \sum_{i=n/2}^n \underbrace{\lg \frac{n}{2}}_{=\lg n - 1} \\
 &= \frac{n}{2} \lg n - \frac{n}{2} = \Omega(n \lg n)
 \end{aligned}$$

- in fact $\lg n! = n \lg n - O(n)$ via [Sterling's Formula](#):

$$n! \sim \sqrt{2\pi n} \left(\frac{n}{e}\right)^n \Rightarrow \lg n! \sim n \lg n - \underbrace{(\lg e)n + \frac{1}{2} \lg n + \frac{1}{2} \lg(2\pi)}_{O(n)}$$

Further algos won't be based on comparison model

Linear-time Sorting (integer sorting)

If n keys are integers ([fitting in a word](#)) $\in 0, 1, \dots, k-1$, can do more than compare them

- \Rightarrow lower bounds don't apply
- if $k = n^{O(1)}$, can sort in $O(n)$ time $O(k)$ always. Also $O(n)$ if $k=n^{O(1)}$
[OPEN](#): $O(n)$ time possible for all k ?

Counting Sort Makes no comparisons

L = array of k empty lists	} $O(k)$
— linked or Python lists	
for j in range n :	
$L[\text{key}(A[j])].\text{append}(A[j])$ $\rightarrow O(1)$	} $O(n)$
random access using integer key	
output = []	} $O(\sum_i (1 + L[i])) = O(k + n)$
for i in range k :	
output.extend($L[i]$)	

Time: $\Theta(n + k)$ — also $\Theta(n + k)$ space

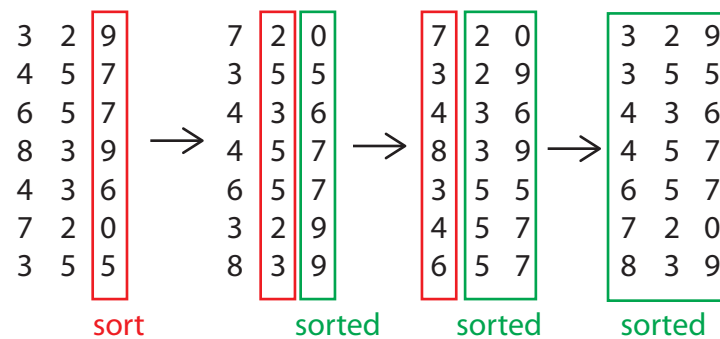
Intuition: Count key occurrences using RAM output <count> copies of each key in order ... but item is more than just a key

CLRS has cooler implementation of counting sort with counters, no lists — but time bound is the same

Radix Sort

- imagine each integer in base b
 $\implies d = \log_b k$ digits $\in \{0, 1, \dots, b - 1\}$ d represents the number of digits
- sort (all n items) by least significant digit \rightarrow can extract in $O(1)$ time
- ...
- sort by most significant digit \rightarrow can extract in $O(1)$ time
sort must be stable: preserve relative order of items with the same key
 \implies don't mess up previous sorting

For example:



- use counting sort for digit sort
 - $\implies \Theta(n + b)$ per digit
 - $\implies \Theta((n + b)d) = \Theta((n + b) \log_b k)$ total time
 - minimized when $b = n$
 - $\implies \Theta(n \log_n k)$
 - $= O(nc)$ if $k \leq n^c$

MIT OpenCourseWare
<http://ocw.mit.edu>

6.006 Introduction to Algorithms
Fall 2011

For information about citing these materials or our Terms of Use, visit: <http://ocw.mit.edu/terms>.